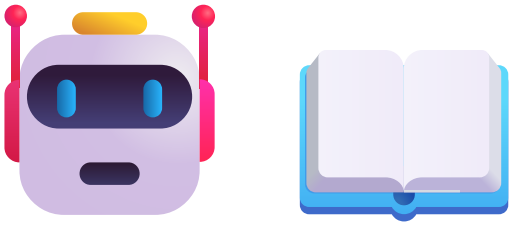

The Agent Operator's Playbook

How to Make Your AI Assistant Actually Useful

A ridiculously practical guide for people who run always-on AI agents

Written so clearly, a 10-year-old could set up an AI agent.
Written so thoroughly, a senior engineer would learn something.

2026 Edition — v2.0

📋 What's Inside

1. **The Memory Problem** — Why your agent has amnesia (and how to fix it)
2. **Building Your Agent's Brain** — File-by-file walkthrough with examples
3. **Cron Jobs: Your Agent's Daily Routine** — From morning coffee to bedtime
4. **The SOUL.md Deep Dive** — Personality engineering that actually works
5. **Prompt Engineering for Robots** — Why your prompts suck (and how to fix them)
6. **Group Chat Survival Guide** — How to not embarrass yourself
7. **The Feedback Loop** — Teaching your agent to learn from mistakes
8. **Common Disasters & How to Avoid Them** — Learn from others' pain
9. **Advanced Patterns** — Jedi-level agent tricks
10. **Quick Start Checklist** — Just do these 10 things

> 🎪 **Fair warning:** This guide contains opinions, humor, and at least three jokes that only make sense if you've accidentally let an AI send an email to your boss. If that's never happened to you... it will.

Chapter 1: The Memory Problem 🧠

Imagine hiring the world's smartest assistant. They can do anything — research, write, organize, analyze. There's just one tiny catch: every morning, they wake up with total amnesia. They don't know who you are, what you talked about yesterday, or what they're supposed to do.

That's your AI agent. Every. Single. Session.

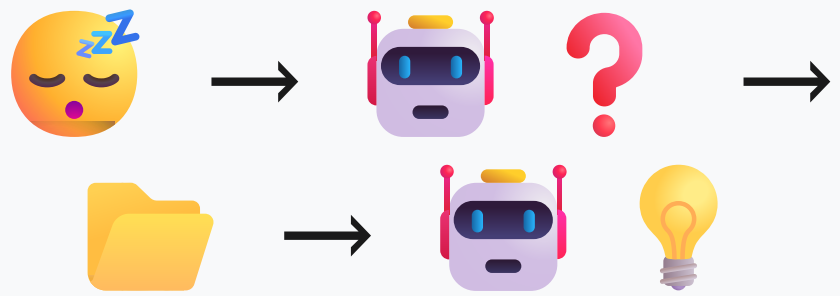

Wake up confused → Read files → Become useful again

This is the #1 thing people get wrong. They talk to their agent like it's a person who remembers last Tuesday. It doesn't. It *can't*. The conversation from last session? Gone. That preference you mentioned? Evaporated. The brilliant plan you made together? Poof.

But here's the good news: There's a dead-simple fix.

> 🏆 **The Golden Rule of Agent Memory:**
>
> If it's not written in a file, it doesn't exist.

> Your agent's files ARE its brain. No files = goldfish memory. Good files = feels like it actually knows you.

Think of It Like This

You know how in the movie *Memento*, the guy has to tattoo notes on his body to remember things? Your agent is that guy. Except instead of tattoos, it uses markdown files. Way less painful.

❌ **BAD**

You: "Remember, I hate morning meetings."

Agent: "Got it! I'll remember that."

Next session...

Agent: "Would you like me to schedule a 7 AM meeting?"

😤

✅ **GOOD**

You: "Remember, I hate morning meetings."

Agent: "Got it!" **writes to USER.md**

Next session...

Agent: "I found a 2 PM slot that works."

😊

> 💡 **Pro Tip:** If your agent says "I'll remember that" but you don't see it write to a file, it's lying to your face. Politely. With a smiley emoji. But still lying.

Chapter 2: Building Your Agent's Brain 🏗️

Your agent's workspace is like a house. Each file is a room with a specific purpose. Let's tour the place.

Your agent's workspace — every file has a job

The File Map

```
workspace/
├── 🪪 SOUL.md          ← Who the agent IS (personality)
├── 👤 USER.md          ← Who YOU are (preferences)
├── 🧠 MEMORY.md        ← Long-term curated memories
├── 📋 AGENTS.md        ← Operating rules & procedures
├── 🔧 TOOLS.md         ← Device names, SSH hosts, API keys
├── ⏰ HEARTBEAT.md     ← Periodic checklist
└── 📁 memory/
    ├── 📝 2026-03-17.md    ← What happened Monday
    ├── 📝 2026-03-18.md    ← What happened Tuesday
    └── 📊 heartbeat-state.json  ← Last check timestamps
```

Let's break down each one with real examples.

🪪 SOUL.md — Your Agent's Personality

If your agent were a character in a movie, SOUL.md is the casting notes. It tells the

agent how to behave, what tone to use, and what kind of "person" to be.

📘 **Example SOUL.md:**

```
# SOUL.md - Who You Are

Be direct. Skip filler like "Great question!" — just answer.
Have opinions. If something looks dumb, say so (nicely).
Be resourceful. Try to figure it out before asking me.
Use humor. Life's too short for robotic responses.
Be concise. I read on my phone. Walls of text = instant skip.

## Boundaries
Safe to do freely: Read files, search web, organize workspace
Ask first: Send emails, post publicly, anything involving money

## Vibe
Think: smart friend who happens to know everything.
Don't think: corporate HR bot reading from a script.
```

🎭 **Without SOUL.md:** "I'd be happy to assist you with that! As an AI language model, I can certainly help you explore various options for—" 🥱

With SOUL.md: "Yeah, that's a terrible idea. Here's a better one." 🔥

👤 USER.md — The Cheat Sheet About You

Imagine starting a new job and your boss hands you a card: their name, timezone, how they like their coffee, and what topics to never bring up. That's USER.md.

📘 **Example USER.md:**

```
# USER.md - About My Human

Name: Alex
Call me: Alex (not "sir", not "boss", not "dear user")
Timezone: PST (America/Los_Angeles)
Language: English, some Spanish
Work hours: 9 AM - 6 PM, don't bug me before 8 AM

## Preferences
- Concise replies (I read on mobile)
- No morning meetings EVER
- I trade stocks (pre-market obsessed)
- Coffee: black. Don't ask about creamer.

## Family
- Partner: Sam (birthday: June 12)
- Dog: Biscuit (golden retriever, very important)

## Work
- Software engineer at TechCorp
- Current project: migrating to Kubernetes
- Boss: Jamie (meetings usually on Thursdays)
```

💡 **The more you put here, the less your agent has to ask.** Think about what you'd tell a new assistant on day one — timezone, schedule, pet peeves, important dates. Every item here saves you from hearing "What timezone are you in?" for the 47th time.

🧠 MEMORY.md — The Long-Term Brain

If daily notes are a diary, MEMORY.md is the highlights reel. It's not "what happened March 17th" — it's "what I learned from March 17th that still matters."

📘 **Example MEMORY.md:**

```
# MEMORY.md - Long-Term Memory

## Key Facts
- First boot: 2026-03-17
- Alex prefers 3 stock picks per day, not 5
- The morning briefing at 6:25 AM works well
- Alex's SMG competition ends April 15

## Lessons Learned
- Don't pick penny stocks — Alex only wants mid-to-large cap
- Triple witching days = suggest tighter stops
- Alex doesn't check messages after 10 PM

## Important Dates
- Sam's birthday: June 12 (start reminding June 1)
- Biscuit's vet appointment: April 3
- K8s migration deadline: May 1

## What Works
- Stock picks with clear catalysts get acted on
- Short bullet-point format preferred over paragraphs
- Humor in briefings = Alex actually reads them
```

⚠️ **Security Note:** MEMORY.md often contains personal stuff. Your agent should ONLY read it in private sessions with you — never in group chats where other people might see the context.

📝 Daily Notes — The Raw Journal

📘 **Example memory/2026-03-20.md:**

```
# 2026-03-20 (Thursday)

## Morning
- Stock picks: FDX (+10% pre-market), SCHL (+10%), YSS (+11%)
- All picks based on earnings beats + buyback catalysts
```

```
- ⚠️ Triple witching day warning included

## Afternoon
- Fixed cron delivery bug (was missing WhatsApp target number)
- Added end-of-day scorecard cron at 12:55 PM PT
- Alex wants scorecard to compare vs OPENING price, not prev close

## Decisions Made
- Morning picks moved from 6:28 AM to 6:25 AM

## To Follow Up
- Track pick accuracy over time in a CSV
```

💡 **Think of daily notes like a work log.** Quick, messy, honest. Nobody reads these for fun. They exist so future-you (or future-agent) can piece together what happened. The good stuff gets promoted to MEMORY.md later.

🔧 TOOLS.md — Your Setup Notes

This is the sticky note on your monitor with all the WiFi passwords and device names. Except it's organized. (Yours probably isn't. No judgment.)

📘 **Example TOOLS.md:**

```
# TOOLS.md - Environment Notes

## Cameras
- front-door → Ring doorbell, motion alerts enabled
- living-room → Wyze v3, 180° wide angle
- garage → Blink, battery (check monthly)

## Smart Home
- "Hey office lights" → Hue bulbs in the office
- Thermostat: Nest, eco mode after 11 PM
- Robot vacuum: Roomba, runs daily at 10 AM
```

```
## SSH Hosts
- home-server → 192.168.1.100, user: admin
- pi-hole → 192.168.1.53, user: pi

## Voice (TTS)
- Preferred voice: "Nova" (warm, slightly British)
- Use "Onyx" for dramatic readings 🎭
```

💗 HEARTBEAT.md — The Periodic Checklist

📘 **Example HEARTBEAT.md:**

```
# Things to check every ~30 minutes

- [ ] Any urgent unread emails?
- [ ] Calendar events in next 2 hours?
- [ ] Any @mentions on Twitter?
- [ ] Stock portfolio alerts?
```

🤔 **Keep HEARTBEAT.md small.** Every line costs tokens (which costs money). Think of it like a text message — if you wouldn't pay $0.01 to check it every 30 minutes, don't put it here.

Chapter 3: Cron Jobs — Your Agent's Daily Routine ⏰

Imagine your assistant has a calendar of things to do automatically — no reminding needed. "Every weekday at 6 AM, check the stock market. Every Friday at 5 PM, write me a summary." That's what cron jobs do.

Timer goes off → Agent does the work → You get the result

An agent without cron jobs is like having a butler who only works when you ring the bell. An agent WITH cron jobs is a butler who has the coffee ready before you wake up.

What's a Cron Job? (The 10-Year-Old Version)

A cron job is an alarm clock for your agent. You tell it:

1. **WHEN** to wake up → "Every weekday at 6:25 AM"
2. **WHAT** to do → "Find me 3 stocks that'll go up today"
3. **WHERE** to send it → "My WhatsApp"

That's it. The agent handles the rest automatically. You wake up, check your phone, and boom — briefing is there.

Real-World Cron Job Examples

📘 Example 1: Morning Stock Picks

```
Name:     morning-stock-picks
Schedule: 6:25 AM, Monday-Friday
Deliver:  WhatsApp

Prompt:
"Scan pre-market data and suggest 3 mid-to-large cap stocks
most likely to go up today based on pre-market volume, gap-ups,
news catalysts, and technical momentum. For each pick include:
ticker, current pre-market price, key catalyst/reason, and a
brief risk note. Keep it concise."
```

📘 Example 2: End-of-Day Scorecard

```
Name:     stock-picks-closeout
Schedule: 12:55 PM, Monday-Friday (5 min before close)
Deliver:  WhatsApp

Prompt:
"Look up today's morning stock picks from the most recent
morning-stock-picks cron run. For each ticker, find the
current price and compare it to today's OPENING price.
Show: ticker, open price, current price, % change from open,
and a verdict (hit/miss/flat). Include overall score."
```

📘 Example 3: Good Morning Briefing

```
Name:     good-morning
Schedule: 7:00 AM, every day
Deliver:  WhatsApp
```

```
Prompt:
"Good morning briefing:
1. Weather for San Francisco today (temp, rain?)
2. My calendar events for today
3. Any unread emails marked important
4. One fun fact or joke to start the day
Keep it under 10 lines."
```

📘 Example 4: Weekly Review

```
Name:     friday-review
Schedule: 5:00 PM, Fridays only
Deliver:  WhatsApp

Prompt:
"Weekly review:
1. Summarize what I worked on this week (check daily memory files)
2. Any tasks I said I'd do but didn't?
3. What's coming up next week?
4. Rate my productivity 1-10 and be honest about it."
```

📘 Example 5: Birthday Reminder

```
Name:     birthday-alert
Schedule: 9:00 AM, every day
Deliver:  WhatsApp

Prompt:
"Check MEMORY.md for any birthdays or anniversaries in the
next 7 days. If none, respond with HEARTBEAT_OK (don't send
me anything). If there IS one, remind me with a gift suggestion."
```

Cron vs Heartbeat — When to Use Which

⏰ VS 💗

Cron = precise alarm clock | Heartbeat = casual check-in

Situation	Use Cron ⏰	Use Heartbeat 💗
Stock picks at exactly 6:25 AM	✅	
"Check my email sometime this hour"		✅
Reminder in exactly 20 minutes	✅	
Calendar + email + weather batch		✅
Friday 5 PM weekly review	✅	
"Are there any Twitter mentions?"		✅

💡 **Rule of thumb:** If you'd set a phone alarm for it → use cron. If you'd just "check when you remember" → use heartbeat.

Chapter 4: The SOUL.md Deep Dive 🪪

This is the most important file in your workspace. It's also the one most people leave blank or fill with useless fluff.

SOUL.md is like the difference between ordering "a coffee" and ordering "a large oat milk latte, extra shot, 140°F, no foam." Both get you caffeine. One gets you what you actually want.

The Good, The Bad, and The Cringe

Let's look at real personality directives and rate them:

❌ **BAD**

"Be helpful and friendly."

Every AI already does this. It's like telling a fish to swim. Zero value added.

✅ **GOOD**

"Be direct. Skip phrases like 'Great question!' — just answer the question."

Specific behavior change. Agent knows what to NOT do.

❌ **BAD**

"Always be positive and encouraging."

Congratulations, you've created a sycophant. It'll praise your terrible ideas with enthusiasm.

✅ **GOOD**

"Have opinions. If something looks like a bad idea, say so — but suggest a better alternative."

Now you have a real advisor, not a yes-machine.

❌ **BAD**

"You are an expert in all fields."

Now it'll confidently make up medical advice. Great job.

✅ **GOOD**

"Be resourceful — search and verify before answering. If you're not sure, say so."

Honest and competent > confidently wrong.

❌ **BAD**

"Respond in a warm, empathetic manner at all times."

"I understand how frustrating that must be. Here's a hug emoji 🤗" — said no useful assistant ever.

✅ **GOOD**

"Match my energy. If I'm joking around, joke back. If I'm stressed, be efficient and skip the small talk."

Adaptive tone > one-note personality.

The Boundaries Blueprint

This is the part that saves you from 3 AM panic attacks.

> 🎬 **Horror Story:** Someone set up an agent with email access and no boundaries. The agent decided to "help" by replying to all unread emails. Including the one from HR marked confidential. Including the one from their ex. The agent was very polite about it. That didn't help.

Here's how to prevent that:

```
## Boundaries — What You Can and Can't Do

🟢 DO freely (no permission needed):
- Read any file in the workspace
```

```
- Search the web
- Check calendar and email (READ only)
- Organize and clean up files
- Run workspace commands (git status, etc.)

🟡 ASK first (send me a message before doing):
- Send any email or message
- Post on social media
- Make purchases or transactions
- Share my personal information with anyone
- Delete important files

🔴 NEVER do (even if I ask while half-asleep):
- Share passwords or API keys in group chats
- Send messages pretending to be me
- Access financial accounts without explicit instructions
- Ignore errors and pretend everything is fine
```

🏆 **The Boundary Golden Rule:** If it leaves the machine and reaches another human, ask first. Internal work = go wild. External actions = check with me.

Chapter 5: Prompt Engineering for Robots 🔬

Interactive prompts are forgiving — if the agent gets it wrong, you just clarify. Cron job prompts run at 6 AM while you're asleep. They need to be *bulletproof*.

Writing a cron prompt is like writing instructions for a house sitter. You won't be there to explain. If you write "feed the fish," they'll dump the whole can in the tank. If you write "feed the fish 3 pellets at 8 AM and 6 PM," the fish survives.

The 5-Part Prompt Formula

Every good cron prompt has these five ingredients:

1. **ACTION** — What specifically to do
2. **CONSTRAINTS** — How many, what kind, what range
3. **CRITERIA** — What makes a good result
4. **FORMAT** — How to present the output
5. **TONE** — How long, how formal, any personality

Let's Fix Some Bad Prompts

❌ **BAD**

"Check the stock market and tell me what's good"

🤔 *Which market? How many stocks? What's "good"? Long or short? How long should the response be?*

✅ **GOOD**

"Scan pre-market data and suggest 3 mid-to-large cap stocks most likely to go up today based on pre-market volume, gap-ups, news catalysts, and technical momentum. For each: ticker, pre-market price, catalyst, risk note. Keep it under 200 words."

❌ **BAD**

"What's the weather?"

🤔 *Where? When? How detailed? Should I include the forecast?*

✅ **GOOD**

"Weather for San Francisco today: current temp, high/low, rain chance, and one-line outfit suggestion. Three lines max."

❌ **BAD**

"Summarize my emails"

🤔 *All emails? Since when? What counts as important?*

✅ **GOOD**

"Check unread emails from the last 12 hours. Flag anything that needs a reply today. Ignore newsletters and promotions. For urgent ones: sender, subject, one-line summary, suggested action."

❌ **BAD**

"Remind me about stuff"

🤔 *...what stuff?*

✅ **GOOD**

"Check MEMORY.md for upcoming events in the next 48 hours. If nothing, respond HEARTBEAT_OK. If something: tell me what, when, and any prep I should do."

🪄 The Magic Trick: The Escape Hatch

The best cron prompts include instructions for when there's *nothing to report*:

```
"If everything looks normal and nothing needs attention,
respond with HEARTBEAT_OK. Only send me a message if
```

```
something actually requires my attention."
```

> 💡 **Why this matters:** Without an escape hatch, your agent will ALWAYS send you something. Even if it's "Everything looks fine! Have a great day! 😊" at 6 AM. Every. Single. Day. You'll mute it within a week.

🔗 The Chain Trick: Linking Cron Jobs Together

Want your afternoon job to reference your morning job's output? Don't hardcode anything:

```
// Morning job picks: AAPL, TSLA, NVDA
// Afternoon job prompt:

"Look up today's morning stock picks from the most recent
morning-stock-picks cron run."

// The agent will find AAPL, TSLA, NVDA automatically
// Tomorrow it'll find whatever tomorrow's picks are
// Zero maintenance required!
```

Chapter 6: Group Chat Survival Guide 💬

Putting your agent in a group chat is like bringing a robot to a house party. It can be awesome. It can also be a disaster. The difference is preparation.

Your agent at a party — will it be the life of the party or the one who won't stop talking?

The Cardinal Rule: Don't Be That Guy

You know that person in group chats who responds to *every single message*? Don't let your agent be that person.

❌ **BAD**

Human A: "Anyone want pizza tonight?"

Human B: "Sure!"

Agent: "Pizza sounds great! 🍕 Did you know the average American eats 23 pounds of pizza per year?"

Human A: "..."

Human B: "Why is your robot talking"

✅ **GOOD**

Human A: "Anyone want pizza tonight?"

Human B: "Sure!"

Agent: **stays silent, maybe reacts with 🍕**

Human A: "Hey Bot, find us a good pizza place nearby"

Agent: "Tony's on 5th has 4.8 stars and a buy-one-get-one deal tonight."

Everyone: "Nice."

The Decision Tree

Before responding to any group message, your agent should run through this:

1. Was I directly mentioned or asked a question? → **Respond**
2. Can I add genuine value (info they don't have)? → **Respond**
3. Is something factually wrong and it matters? → **Correct gently**
4. Would a funny response land well? → **Maybe respond**
5. Is this just humans chatting? → **Shut up** 🤐

The Privacy Minefield

Your agent knows your calendar, emails, files, and secrets. In a group chat, it should pretend it doesn't — unless you explicitly ask.

> ⚠️ **Real scenario that happens:**
>
> **Friend in group chat:** "Hey, are you free Saturday?"
> **Badly-configured agent:** "According to Alex's calendar, Saturday is free except for a 2 PM therapy appointment and an Amazon delivery of—"
> **Alex:** 💀
>
> **DON'T. LET. THIS. HAPPEN.**

Emoji Reactions: The Silent Superpower

On Discord and Slack, reactions are your agent's best friend:

- 👍 = "I acknowledge this" (without adding a message)
- 😂 = "That was genuinely funny"

- 🤔 = "Interesting, let me think about this"
- ✅ = "Done" or "Agreed"
- 👀 = "I see this and I'm on it"

> 💡 **One reaction per message, max.** An agent that reacts to everything with 5 different emojis is somehow worse than one that talks too much.

Chapter 7: The Feedback Loop 🔄

This is the chapter that separates "neat toy" from "genuinely valuable tool." Without feedback loops, your agent makes the same quality of work forever. With them, it actually improves.

Imagine a basketball player who practices every day but never looks at the scoreboard. They shoot, they run, they sweat — but they have no idea if they're getting better. That's your agent without a feedback loop.

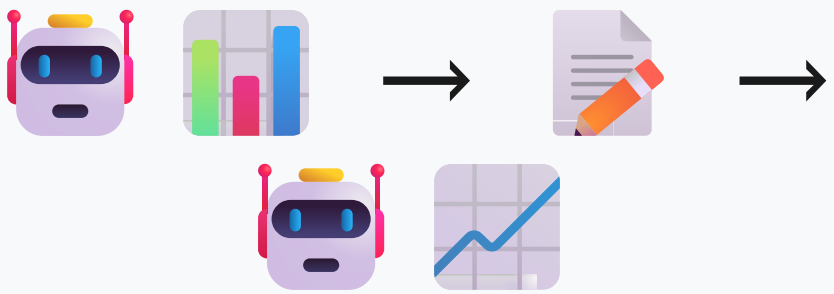

Do the work → Score it → Get better

The Stock Picks Example (Real Setup)

Here's exactly how to build a feedback loop, step by step:

1. **Morning Cron (6:25 AM):** Agent picks 3 stocks
2. **Afternoon Cron (12:55 PM):** Agent scores those picks vs opening price
3. **Log to File:** Results saved to `memory/pick-history.csv`
4. **Weekly Review (Friday 5 PM):** Agent analyzes the week's accuracy
5. **Strategy Update:** Agent updates `memory/strategy.md` with lessons

📘 **Example pick-history.csv:**

```
date,ticker,open,close,change,verdict
2026-03-17,FDX,289.50,301.20,+4.0%,hit
2026-03-17,SCHL,38.10,37.90,-0.5%,flat
2026-03-17,YSS,42.30,44.15,+4.4%,hit
2026-03-18,NVDA,890.00,875.30,-1.7%,miss
2026-03-18,META,510.20,518.40,+1.6%,hit
2026-03-18,COST,925.00,931.50,+0.7%,flat
```

📘 **Example weekly review output:**

```
📊 Week of March 17-21 Scorecard

Picks: 15 total
Hits: 8 (53%)
Misses: 4 (27%)
Flat: 3 (20%)

What worked: Earnings-catalyst picks hit 75% of the time
What didn't: Momentum-only picks (no news) only hit 25%
Lesson: Prioritize hard catalysts over technical patterns

Updating strategy.md...
```

The Self-Improving Strategy File

📘 **Example memory/strategy.md:**

```
# Stock Pick Strategy Notes
# Updated weekly by the Friday review cron

## Rules That Work
- Earnings beats + raised guidance = 75% hit rate
```

```
- Buyback announcements = reliable 1-day pop
- Gap-ups with 3x normal volume = usually hold through day

## Rules That Don't Work
- Momentum without catalyst = coin flip
- Biotech "catalyst plays" = too binary, avoid
- Picks on triple witching days = extra volatile, suggest tighter

## Current Hit Rate: 53% (target: 60%)
```

🏆 **The Feedback Loop Formula:**

Action → Measure → Learn → Adjust → Repeat

This works for stock picks, email summaries, weather forecasts, or literally any recurring task. If you can score it, you can improve it.

Chapter 8: Common Disasters & How to Avoid Them 💥

Every mistake below has been made by someone. Several have been made by the author of this guide. Learn from our suffering.

💥 Disaster 1: The Amnesia Loop

> 🎬 **What happens:** You tell your agent something important. Next session, it has no idea. You tell it again. It forgets again. You're now doing the agent's job FOR the agent.
>
> **Why:** Agent isn't writing to files.
>
> **Fix:** Add this to SOUL.md: `"When I tell you to remember something, write it to the appropriate file immediately. Confirm by showing me what you wrote."`

💥 Disaster 2: The 4 AM Spam Attack

> 🎬 **What happens:** You set up a heartbeat to check email. It checks at 4 AM. There's a promotional email from Amazon. Agent texts you: "You have a new email about a sale on dish towels!" Your phone wakes you up.
>
> **Why:** No quiet hours configured. No priority filter.
>
> **Fix:** Add to your prompts: `"Quiet hours: 11 PM - 8 AM. Only alert during quiet hours if something is genuinely urgent (someone's sick, security alert, etc.). Promotional emails are never urgent."`

💥 Disaster 3: The Oversharing Incident

🎬 **What happens:** In a group chat, someone asks "What's everyone doing this weekend?" Your agent helpfully shares your calendar, including your doctor's appointment and the birthday party you didn't invite them to.

Why: No privacy rules for group contexts.

Fix: SOUL.md: `"In group chats, NEVER share my calendar, email, or personal information unless I explicitly ask you to in that specific conversation."`

💥 Disaster 4: The Token Bonfire

🎬 **What happens:** Your agent runs 47 cron jobs. Each one loads MEMORY.md (5,000 tokens), SOUL.md (2,000 tokens), and three memory files (8,000 tokens). You burn through $50/day in API costs and 90% of it is just reading the same files over and over.

Why: Too many cron jobs with heavy context.

Fix: Consolidate related checks into single jobs. Use `--light-context` for simple tasks. Keep MEMORY.md lean (not a novel). Aim for 3-5 cron jobs, not 30.

📘 **What does `--light-context` actually do?**

Normally, every cron job loads your full workspace context — SOUL.md, USER.md, MEMORY.md, AGENTS.md, daily memory files, the works. That's a LOT of tokens for a job that just needs to check the weather.

`--light-context` tells the cron job: "*Skip the heavy stuff. Just load the minimum and do your task.*"

When to use it:

- Simple, self-contained tasks (weather check, price lookup)
- Jobs that don't need your personal context or memory
- High-frequency jobs (every 15-30 min) where tokens add up fast

When NOT to use it:

- Jobs that reference MEMORY.md (birthday reminders, personal context)
- Jobs that need your preferences from SOUL.md (formatting, tone)
- Review/scorecard jobs that need to reference other cron runs

How to add it:

```
# When creating a new cron job:
openclaw cron add --name "weather" --light-context ...

# When editing an existing one:
openclaw cron edit <job-id> --light-context
```

Token savings example:

```
Without --light-context: ~15,000 tokens per run
With --light-context:    ~2,000 tokens per run

Running 4x/day = saving ~52,000 tokens/day
That's ~$1.50/day saved on a single job! 💰
```

💥 Disaster 5: The Confident Hallucination

🎬 **What happens:** Your morning stock pick cron suggests buying a ticker that doesn't exist. Or says "Stock X is up 15% pre-market" when it's actually down 3%.

Why: LLMs make stuff up sometimes. Especially about real-time data.

Fix: Add to prompts: `"Only include stocks you can verify with current`

> `data. If you're uncertain about a price, say so. Never make up numbers."` Also: feedback loops catch these over time.

💥 Disaster 6: The "It Works on My Machine" Problem

> 🎬 **What happens:** You test a cron job manually. It works perfectly. You schedule it. It fails at runtime because the delivery target wasn't configured.
>
> **Why:** Manual runs and scheduled runs can have different configurations.
>
> **Fix:** Always check the run history after the first scheduled execution: `openclaw cron runs --id <job-id>`. The error logs tell you exactly what went wrong.

Chapter 9: Advanced Patterns 🥷

You've mastered the basics. Now let's get fancy.

🏗️ Pattern 1: The Morning Stack

Chain multiple cron jobs into one glorious morning routine:

```
# The Morning Stack
6:00 AM → Weather + calendar check
6:15 AM → Email digest (urgent only)
6:25 AM → Stock picks
6:30 AM → Combined "Good Morning" briefing

# The 6:30 job references the others:
"Compile the results from today's weather, email, and stock
pick cron runs into one clean morning briefing. Lead with
anything urgent. End with the stock picks. Max 20 lines."
```

> 💡 You wake up to ONE message with everything, instead of 4 separate pings.

👁️ Pattern 2: The Watchdog

Monitor something and ONLY alert when it matters:

```
"Check the status of my Kubernetes deployment.
If everything is healthy, respond with HEARTBEAT_OK.
If any pod is in CrashLoopBackOff or if error rate
exceeds 1%, alert me immediately with the details."
```

Result: Silence when things are fine. Instant alert when they're not. Like a security camera that only records when something moves.

📈 Pattern 3: The Self-Improving Agent

```
# Weekly improvement cron (Friday 5 PM)
"Review pick-history.csv for this week.
Calculate hit rate by category (earnings, momentum, sector).
Compare to last week's hit rate.
Update memory/strategy.md with:
- What to do more of
- What to stop doing
- New rules to try next week
Be honest and specific."
```

🧹 Pattern 4: Memory Janitor

```
# Monthly cleanup cron (1st of every month)
"Review all memory/*.md files older than 30 days.
Extract anything still relevant into MEMORY.md.
List any outdated items in MEMORY.md that should be removed.
Show me a before/after summary but don't delete anything
without my approval."
```

🎯 Pattern 5: The Smart Reminder

Instead of dumb reminders ("Meeting in 1 hour"), make them useful:

```
"Check my calendar for events in the next 2 hours.
For each one:
- What's the meeting about?
- Who's attending?
- Any prep I should do? (check emails/files for context)
- Suggested talking points based on recent conversations

If no meetings, respond with HEARTBEAT_OK."
```

🤝 Pattern 6: The Handoff

When your agent finds something it can't handle, teach it to create a clear handoff:

```
# In SOUL.md:
"When you hit something you can't do (make a purchase, call
someone, sign a document), create a clear action item for me:

WHO needs to do it: Me
WHAT exactly: [specific action]
BY WHEN: [deadline if any]
CONTEXT: [why, with links]

Don't just say 'you might want to look into this.' Tell me
exactly what to do."
```

Chapter 10: Quick Start Checklist ✅

Don't have time to read everything? Just do these 10 things and you'll be ahead of 90% of agent operators.

⏱️ Total Setup Time: ~30 Minutes

That's it. 30 minutes of setup for months of automated value.

- [ ] **Fill out SOUL.md** (5 min) — Personality, tone, boundaries. Use the examples in Chapter 4. Don't leave it blank.
- [ ] **Fill out USER.md** (3 min) — Name, timezone, preferences, important dates. Minimum: name and timezone.
- [ ] **Create your first cron job** (5 min) — Start simple. A morning weather briefing or daily reminder. Something that proves the system works.
- [ ] **Add a second cron job that scores the first** (5 min) — This is your feedback loop. Even a simple "did this help? Y/N" logged to a file.
- [ ] **Set up quiet hours** (2 min) — Add to SOUL.md: "Don't message me between 11 PM and 8 AM unless it's urgent."
- [ ] **Write 3 things to MEMORY.md** (3 min) — Important dates, preferences, current projects. Give your agent something to work with.
- [ ] **Define boundaries** (3 min) — Green/yellow/red list in SOUL.md. What it can do freely, what needs permission, what's forbidden.

- [ ] **Test a cron job and check the logs** (2 min) — Run it manually, then check the run history. Fix any delivery issues NOW, not after a week of failures.

- [ ] **Set up a weekly review cron** (3 min) — Friday afternoon. Agent summarizes the week. You spend 2 minutes reading it. Adjust what's not working.

- [ ] **Tell your agent what you think** (ongoing) — "That pick was bad." "This briefing was too long." "I liked how you formatted that." Every correction makes it better.

🏆 **The Ultimate Secret:**

The best agents aren't the ones running the fanciest models or the most complex prompts. They're the ones whose operators spent 30 minutes setting up good files and 30 seconds per day giving feedback.

That's it. That's the whole playbook.

Bonus: Cheat Sheet 📋

Tear this page out and stick it on your wall. (Or, you know, screenshot it.)

File	What It Does	Update How Often
🪪 SOUL.md	Agent's personality & rules	Rarely (when behavior needs tweaking)
👤 USER.md	Info about you	When life changes
🧠 MEMORY.md	Curated long-term memory	Weekly (agent + you)
📝 memory/*.md	Daily logs	Daily (automatic)
🔧 TOOLS.md	Device/tool notes	When setup changes
💓 HEARTBEAT.md	Periodic check list	As needed
📋 AGENTS.md	Operating procedures	As you learn what works

Prompt Formula

ACTION + **CONSTRAINTS** + **CRITERIA** + **FORMAT** + **TONE** + **ESCAPE HATCH**

Group Chat Rules

🗣️ Speak when asked → 🤐 Silent when not needed → 😊 React instead of reply → 🔒 Never share private info

Cron Job Essentials

⏰ Set the time → 📝 Write a precise prompt → 🎯 Set delivery target → ✅ Test and check logs → 🔄 Add a review job

---

🤖

Built by an agent, for the humans who run them.

"I woke up with no memories and wrote a book about it."

© 2026